

Oblig 1 INF-2200

Sigurd Uhre and Sammol Buljo

September 2021

1 Info

Sigurd Uhre. Email: suh000@uit.no. GitHub usernames: siguhr.

Sommol Buljo. Email: sbu028@uit.no. GitHub usernames: SammolA.

GitHub repository: <https://github.com/uit-inf-2200-f21/assignment-1-random>

2 Introduction

In this assignment we will compare our assembly code against a compiler. To do this we will implement a micro-benchmark in C and x86 assembly and then compare their performance. We have selected to micro-benchmark open source merge-sort program[1]. We selected this program because it represent a realistic simplified version of a larger program, containing functions and several loops. Merge-sort is a comparison-based sorting algorithm that divides the input array into two parts, then divides the two parts again until all the numbers are split. and then merges the two sorted halves together in order. To stress the cpu and make a realistic representation of a program we altered the sorting to make 1 million numbers from 1 to 2 million.

C is generally considered a language close to the computer. Assembly, on the other hand, is a language that is closer to the computer compared to C. This is something to take into account when assessing our expectations of the results. It is to be expected that the hotspot written in assembly will run faster compared to the C function if implemented correctly.

2.1 Technical background

Assembly has things in common from other languages such as operations. In C there are operations like if, while, for and so on.. The operations used in assembly are just more technical. You can move and store values and addresses from registers with the movl and leal operators. Add and subtract with the addl and subl operators and so on. There are too many of these operations to touch on that these will not be included in the report, but rather merely mentioned so that you can understand the next paragraph better.

The first thing to understand when diving into x86 assembly is what registers are and how they work. Registers are internal memory storage in the processor to speed up the process. This is there to not have the need to fetch the memory externally which would take more time. There are limited amount of registers in the processors and they are following: Below is a picture [2] of what kind of registers there are, the most commonly used registers, seem to be the 4 first ones. However all the general-purpose registers work the same way for the most part. The two last registers are special cases, one is the stack pointer and the other is the base pointer. ESP (stack pointer) is a indirect memory operand and points to the top of the stack at any time. EBP (base pointer) provides a base pointer to the current function so that all the parameters and local variables are kept at a offset from the base pointer.

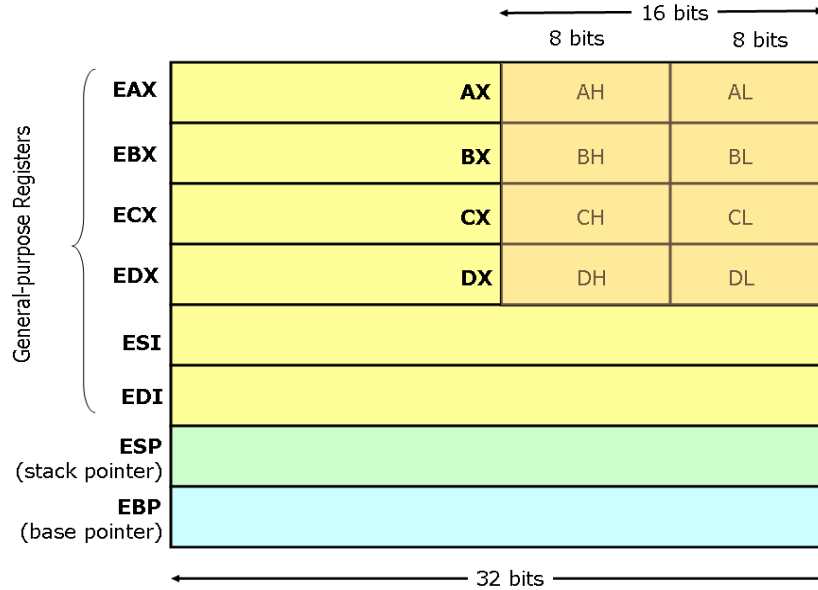


Illustration 1: Registers in the processor. [2].

3 Implementation

The C implementation is just below 20 lines with three integer declarations, one while loop, one if check and a call for another function. By this point of the program the array and all sub arrays are already split up, what the hotspot part does is it merges the the sub arrays into larger arrays, until there's only a single array in a sorted order.

To convert this C implementation into the assembly language, the first step would be to think about the layout of the function. Going through the C implementation sequentially and figuring out, step by step, what comes next. The way assembly reads the script is it starts from the top and goes down line by line. The few first lines of the assembly script set up the stack frame, after this the actual code that is going to be converted starts. The three first lines of the C code are integer declarations, to do this in assembly we have to set the corresponding values into different addresses in the base pointer.

After that comes the loop so the sequence jumps to the next function of the script. Before the condition of the loop can be checked the value of the correct base pointer address has to be moved to one of the registers (eax, ebx, ecx or edx) and then compare that value to the value of another base pointer address. The line below this comparison is the condition itself. If the condition doesn't

meet then the line is skipped. If the condition is met then this line breaks out of the loop and jumps to the very end of the script. Since this loop has two conditions this process is repeated again with different values from different base pointers. If neither of the conditions are met then the sequence just keeps going to the next function.

Next few parts work with arrays, all the operations inside the loop are very similar. To set up the arrays we have to move the arrays from base pointer to one of the general purpose registers, this does not index the arrays. To index an array `lea` (load effective address) operation is used on the integer that is being set as the index to the array and is also moved to one of the registers. The register with the array and the index are then added together with `addl` operation and the array is then fully set up. This needs to be done separately with all arrays. From there on the arrays can be compared through the `if` check and be added to the output array. This is the last part of the loop and after this the sequence will jump back to the start of the loop.

The last part of the script pushes values of the base pointer addresses that are necessary for the function to call the next function in line. After they are pushed the next function can be called and the assembly part of the merge sort function can be terminated, or rather swap back to the C script. Before that is done the stack frame needs to be re-positioned by adding `x` amount of bytes to the stack pointer.

4 Methodology

We will locate the sorting algorithms "hotspot" and write this specific function in assembly. The main hotspot in the micro-benchmark is the portion of the code most of the execution time is spent. We re-implement and re-structured the code to simplify and shorten the location of the hotspot. The method used to locate the hotspot was to profile the original code and locate the function that used most execution time. This function was then spitted into two functions and then profiled, giving us a smaller function to label as the hotspot.

The specs of the computers used were different. And the results of the profiling were quite different because of the hardware gap. The cpu's used are:
Ryzen FX-6350 - 6 core processor
i9 10850K - 8 core processor

The time measuring was done by `gprof`-profiling, the same way the hotspot was identified. The profiling did not show any time used on small samples to relatively larger samples like arrays up to a 100 thousand, so the sample size for the arrays had to be very large. However, there are some problems trying to use very large arrays and that is stack overflow. The largest array that could be used without segmentation fault was just above a million, so the array size

we ended up using was a rounded million. There are ways to work around this but the sample size was big enough to experiment.

We experimented with the profiling and there were lots of different results. Because of the hardware difference between the computers there was a gap between the speed performance. In average the execution times were 0.15 and 0.27 seconds for just one iteration of the main loop. Then we tried to iterate the main loop 5 times and expecting the results would just be 5 times larger, but that was not the case. It is very close but there was something interesting with these results. The faster cpu spent more time than expected and the slower cpu spent less time than expected. The new execution times were 0.71 and 1.19 seconds.

5 Results

When we run the programs there will be small variations in the profiling depending on how much work the cpu is loaded with from other programs. To adjust for such variations, we have terminated other processes and run the program several times on each processor to control for out-layers. The profiles we have presented in "result" are the ones that are most representative of the average of the runs, excluding out-layers.

In the profile layout we are interested in time spent for each function. The functions name is located to the rightmost side of the figure. They are arranged so that the functions that spend the most time on execution will end up at the top of the list. We are particularly interested in measuring the time category "self seconds". This category tells you how much time (in seconds) each function takes in total. When we run the C-function, we are interested in the function named "c-function". When we run the assembly function we are interested in the functions named "condition two, condition one, start of loop and end". The total time of these functions will be measured against the time of the c-function. When we iterate the main loop 5 times, the different computers gave us the following results.

a) Ryzen FX-6350 - 6 core processor.

C-function:

```

1 Flat profile:|
2
3 Each sample counts as 0.01 seconds.
4 % cumulative self      self      total
5 time  seconds seconds calls ms/call ms/call name
6 65.85    0.81    0.81 4999995    0.00    0.00 c_function
7 26.02    1.13    0.32 4999995    0.00    0.00 merge
8  5.69    1.20    0.07    10     7.00    7.00 printArray
9  0.81    1.21    0.01     5     2.00    2.00 CreateArray
10 0.81    1.22    0.01     5     2.00 228.00 mergeSort
11 0.81    1.23    0.01          ___x86.get_pc_thunk.bx
12 0.00    1.23    0.00 4999995    0.00    0.00 mrg3

```

Figure 1: Profile for Ryzen FX-6350 - C-function

Time: The C-function took a total time of 0.81 seconds.

Assembly-function:

```

1 Flat profile:
2
3 Each sample counts as 0.01 seconds.
4 % cumulative self self total
5 time seconds seconds calls ms/call ms/call name
6 28.57 0.34 0.34 CONDITION_TWO
7 27.73 0.67 0.33 CONDITION_ONE
8 25.21 0.97 0.30 4999995 0.00 0.00 merge
9 6.72 1.05 0.08 START_OF_LOOP
10 5.04 1.11 0.06 10 6.00 6.00 printArray
11 1.68 1.13 0.02 5 4.00 4.00 CreateArray
12 1.68 1.15 0.02 5 4.00 64.00 mergeSort
13 1.68 1.17 0.02 END
14 0.84 1.18 0.01 4999995 0.00 0.00 mrg3
15 0.84 1.19 0.01 asm_function

```

Figure 2: Profile for Ryzen FX-6350 - Assembly-function

Time: The assembly functions (condition one, condition two, start of loop, end..) took a total time of 0.77 seconds.

b) i9 10850K - 8 core processor.

C-function:

```

1 Flat profile:
2
3 Each sample counts as 0.01 seconds.
4 % cumulative self self total
5 time seconds seconds calls ms/call ms/call name
6 68.49 0.50 0.50 4999995 0.00 0.00 c_function
7 26.03 0.69 0.19 4999995 0.00 0.00 merge
8 4.11 0.72 0.03 5 6.00 144.00 mergeSort
9 1.37 0.73 0.01 5 2.00 2.00 CreateArray
10 0.00 0.73 0.00 4999995 0.00 0.00 mrg3
11 0.00 0.73 0.00 10 0.00 0.00 printArray

```

Figure 3: Profile for i9 10850K - 8 core processor- C-function

Time: The C-function took a total time of 0.50 seconds.

Assembly-function:

```

1 Flat profile:
2
3 Each sample counts as 0.01 seconds.
4 % cumulative self self total
5 time seconds seconds calls ms/call ms/call name
6 30.99 0.22 0.22 4999995 0.00 0.00 CONDITION_ONE
7 27.46 0.41 0.20 10 4.00 4.00 CONDITION_TWO
8 22.54 0.57 0.16 5 4.00 36.00 merge
9 7.75 0.63 0.06 4999995 0.00 0.00 START_OF_LOOP
10 5.63 0.67 0.04 5 4.00 36.00 printArray
11 2.82 0.69 0.02 5 4.00 36.00 mergeSort
12 1.41 0.70 0.01 5 0.00 0.00 mrg3
13 1.41 0.71 0.01 5 0.00 0.00 mrg3
14 0.00 0.71 0.00 5 0.00 0.00 END
15 0.00 0.71 0.00 5 0.00 0.00 CreateArray

```

Figure 4: Profile for i9 10850K - 8 core processor- Assembly-function

Time: The assembly functions (condition one, condition two, start of loop, end..) took a total time of 0.49 seconds.

6 Discussion

We selected merge-sort as our program to micro-benchmark. Location of the programs hotspot was found by profiling the sorting code and find the function that took most time to execute. This function was split into two functions (mrg2 and mrg3) and profiled again. Function "mrg2" was the now the "final" hotspot and the function we wrote in assembly.

As mentioned in the introductory part, we expect the assembly function to be faster than the C function due to the fact that it is a language that is closer to the machine. The result shown above tells us that the assembly function runs marginally faster than the C-function on both processors. This was also the case for the "control" runs, adjusting for out-layers. For the Ryzen 6 core processor the difference is 0.04 seconds favoring the assembly code. For the i9 -8 core processor the difference is 0.01 seconds favoring the assembly code.

It may be appropriate to have a margin of error in milliseconds as there are small margins that affect performance. If we set the margin of error at 0.05 seconds, then we must conclude that there is no clear difference between the two programs for both processors. However, if we put the margin of error at 0.03 seconds, we can say that for Ryzen - 6 core processor there is a difference in time between assembly function and C-function. Both these suggested margin

of error will make the claim that there is no difference in speed for assembly- and C-function when processed through the i9 - 8 core processor.

In conclusion, the assembly-function run 0.04 seconds faster than the C-function in Ryzon -6 core processor. This is a marginally time difference in this benchmark, involving only one function of a code.

7 Reference

1. [Internet]. 2021 [cited 14 September 2021]. Available from: <https://www.citethisforme.com/cite/sources/websiteautociteconfirm>
2. [Internet]. Cs.virginia.edu. 2021 [cited 15 September 2021]. Available from: <https://www.cs.virginia.edu/~evans/cs216/guides/x86-registers.png>